



Area Plots, Histograms, and Bar Plots

Estimated time needed: 30 minutes

Objectives

After completing this lab you will be able to:

- Create additional labs namely area plots, histogram and bar charts

Table of Contents

1. Exploring Datasets with "pandas" (#0)
2. Downloading and Prepping Data (#2)
3. Visualizing Data using Matplotlib (#4)
4. Area Plots (#6)
5. Histograms (#8)
6. Bar Charts (#10)

Exploring Datasets with *pandas* and *Matplotlib*

Toolkits: The course heavily relies on **pandas** and **Numpy** for data wrangling, analysis, and visualization. The primary plotting library that we are exploring in the course is **Matplotlib**.

Dataset: Immigration to Canada from 1980 to 2013 - [International migration flows to and from selected countries - The 2015 revision from United Nations](#) website.

The dataset contains annual data on the flows of international migration as recorded by the countries of destination. The data presents both inflows and outflows according to the place of birth, citizenship or place of previous / next residence both for foreigners and nationals. For this lesson, we will focus on the Canadian Immigration data.

Downloading and Prepping Data

Import Primary Modules. The first thing we'll do is import two key data analysis modules: `pandas` and `numpy`.

```
import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library
```

Let's download and import our primary Canadian Immigration dataset using `pandas`' `read_excel()` method. Normally, before we can do that, we would need to download a module which `pandas` requires reading in Excel files. This module was **openpyxl** (formerly **xlrd**). For your convenience, we have pre-installed this module, so you would not have to worry about that. Otherwise, you would need to run the following line of code to install the **openpyxl** module:

```
! pip3 install openpyxl
```

Download the dataset and read it into a `pandas` dataframe.

```
df_can = pd.read_excel(
    "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DV0101EN-SkillsNetwork/Data%20Science/Sheet1.xlsx",
    sheet_name='Canada by Citizenship',
    skipfooter=2,
    skipheader=1,
    usecols='A:Z')
print("Data downloaded and read into a dataframe!")
```

Data downloaded and read into a dataframe!

Let's take a look at the first five items in our dataset.

```
df_can.head()
```

	Type	Coverage	OdName	AREA	AreaName	REG	RegName	DEV	DevName	1980	1981	2004	2005	2006	2007	2008	2009
0	Immigrants	Foreigners	Afghanistan	935	Asia	5501	Southern Asia	902	Developing regions	16	...	2978	3436	3009	2652	2111	1746
1	Immigrants	Foreigners	Albania	908	Europe	925	Southern Europe	901	Developed regions	1	...	1450	1223	856	702	560	716
2	Immigrants	Foreigners	Algeria	903	Africa	912	Northern Africa	902	Developing regions	80	...	3616	3626	4807	3623	4005	5393
3	Immigrants	Foreigners	American Samoa	909	Oceania	957	Polynesia	902	Developing regions	0	...	0	0	1	0	0	0
4	Immigrants	Foreigners	Andorra	908	Europe	925	Southern Europe	901	Developed regions	0	...	0	0	1	1	0	0

5 rows × 17 columns

Let's find out how many entries there are in our dataset.

```
# print the dimensions of the dataframe
print(df_can.shape)
```

(195, 17)

Clean up data. We will make some modifications to the original dataset to make it easier to create our visualizations. Refer to [Introduction to Matplotlib and Line Plots](#) lab for the rational and detailed description of the changes.

1. Clean up the dataset to remove columns that are not informative to us for visualization (eg. Type, AREA, REG).

```
df_can.drop(['AREA', 'REG', 'DEV', 'Type', 'Coverage'], axis=1, inplace=True)
# let's view the first five elements and see how the dataframe has changed
df_can.head()
```

```
df_can.head()
```

	OdName	AreaName	RegName	DevName	1980	1981	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
0	Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	...	2978	3436	3009	2652	2111	1746	1758	2203	2203	
1	Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	...	1450	1223	856	702	560	716	561	539	539	
2	Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	...	3616	3626	4807	3623	4005	5393	4752	4325	4325	
3	American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0
4	Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	...	0	0	1	1	0	0	0	0	0	0

5 rows × 22 columns

Notice how the columns Type, Coverage, AREA, REG, and DEV got removed from the dataframe.

2. Rename some of the columns so that they make sense.

```
df_can.rename(columns={'OdName':'Country', 'AreaName':'Continent', 'RegName':'Region'}, inplace=True)
# let's view the first five elements and see how the dataframe has changed
df_can.head()
```

```
df_can.head()
```

	Country	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
0	Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	...	2978	3436	3009	2652	2111	1746	1758	2203	2203	
1	Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	...	1450	1223	856	702	560	716	561	539	539	
2	Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	...	3616	3626	4807	3623	4005	5393	4752	4325	4325	
3	American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0
4	Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	...	0	0	1	1	0	0	0	0	0	0

5 rows × 22 columns

Notice how the column names now make much more sense, even to an outsider.

For consistency, ensure that all column labels of type string.

```
# let's examine the types of the column labels
all(isinstance(column, str) for column in df_can.columns)
```

True

Notice how the above line of code returned `True` when we tested if all the column labels are of type `string`. So let's change them all to `string` type.

```
df_can.columns = list(map(str, df_can.columns))
# let's check the column labels types now
all(isinstance(column, str) for column in df_can.columns)
```

True

4. Set the country name as index - useful for quickly looking up countries using `.loc` method.

```
df_can.set_index('Country', inplace=True)
# let's view the first five elements and see how the dataframe has changed
df_can.head()
```

```
df_can.head()
```

Country	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	...	2978	3436	3009	2652	2111	1746	1758	2203	2203	
Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	...	1450	1223	856	702	560	716	561	539	539	
Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	...	3616	3626	4807	3623	4005	5393	4752	4325	4325	
American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0
Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	...	0	0	1	1	0	0	0	0	0	0

5 rows × 22 columns

Notice how the column names now serve as indices.

5. Add total column.

```
df_can['Total'] = df_can.sum(axis=1)
# let's view the first five elements and see how the dataframe has changed
df_can.head()
```

```
df_can.head()
```

Country	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	...	2978	3436	3009	2652	2111	1746	1758	2203	2203	2203	17588
Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	...	1450	1223	856	702	560	716	561	539	539	539	539
Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	...	3616	3626	4807	3623	4005	5393	4752	4325	4325	4325	4325
American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0	0
Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	...	0	0	1	1	0	0	0	0	0	0	0

5 rows × 23 columns

Now the dataframe has an extra column that presents the total number of immigrants from each country in the dataset from 1980 - 2013.

So if we print the dimension of the data, we get:

```
print('data dimensions:', df_can.shape)
```

data dimensions: (195, 23)

So now our dataframe has 38 columns instead of 37 columns that we had before.

```
# Finally, let's create a list of years from 1980 - 2013
years = list(map(str, range(1980, 2014)))
```

years

```
years
```

['1980', '1981', '1982', '1983', '1984', '1985', '1986', '1987', '1988', '1989', '1990', '1991', '1992', '1993', '1994', '1995', '1996', '1997', '1998', '1999', '2000', '2001', '2002', '2003', '2004', '2005', '2006', '2007', '2008', '2009', '2010', '2011', '2012', '2013']

Visualizing Data using Matplotlib

Import the `matplotlib` library.

```
# use the inline backend to generate the plots within the browser
%matplotlib inline
import matplotlib.pyplot as plt
```

```
plt.style.use('ggplot') # optional: for ggplot-like style
# check for latest version of Matplotlib
print('Matplotlib version: ', plt.__version__) # >= 2.0.0
```

UsageError: Line magic function '%matplotlib inline' not found.

Area Plots

In the last module, we created a line plot that visualized the top 5 countries that contributed the most immigrants to Canada from 1980 to 2013. With a little modification to the code, we can visualize this plot as a cumulative plot, also known as a **Stacked Line Plot** or **Area Plot**.

```
df_can.sort_values(['Total'], ascending=False, axis=0, inplace=True)
```

```
# get the top 5 entries
df_top5 = df_can.head()
```

```
# transpose the dataframe
df_top5 = df_top5.T
```

```
df_top5.head()
```

```
df_top5.head()
```

	Country	India	China	United Kingdom of Great Britain and Northern Ireland	Philippines	Pakistan	
1980		8880	5123		22045	6051	978
1981		8670	6682		24796	5921	972
1982		8147	3308		20620	5249	1201
1983		7338	1863		10015	4562	900
1984		5704	1527		10170	3801	668

5 rows × 6 columns

```
df_top5 = df_top5.T
df_top5.index = df_top5.index.map(int)
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```

```
df_top5.plot(kind='area', stacked=False, figsize=(20,10))
```



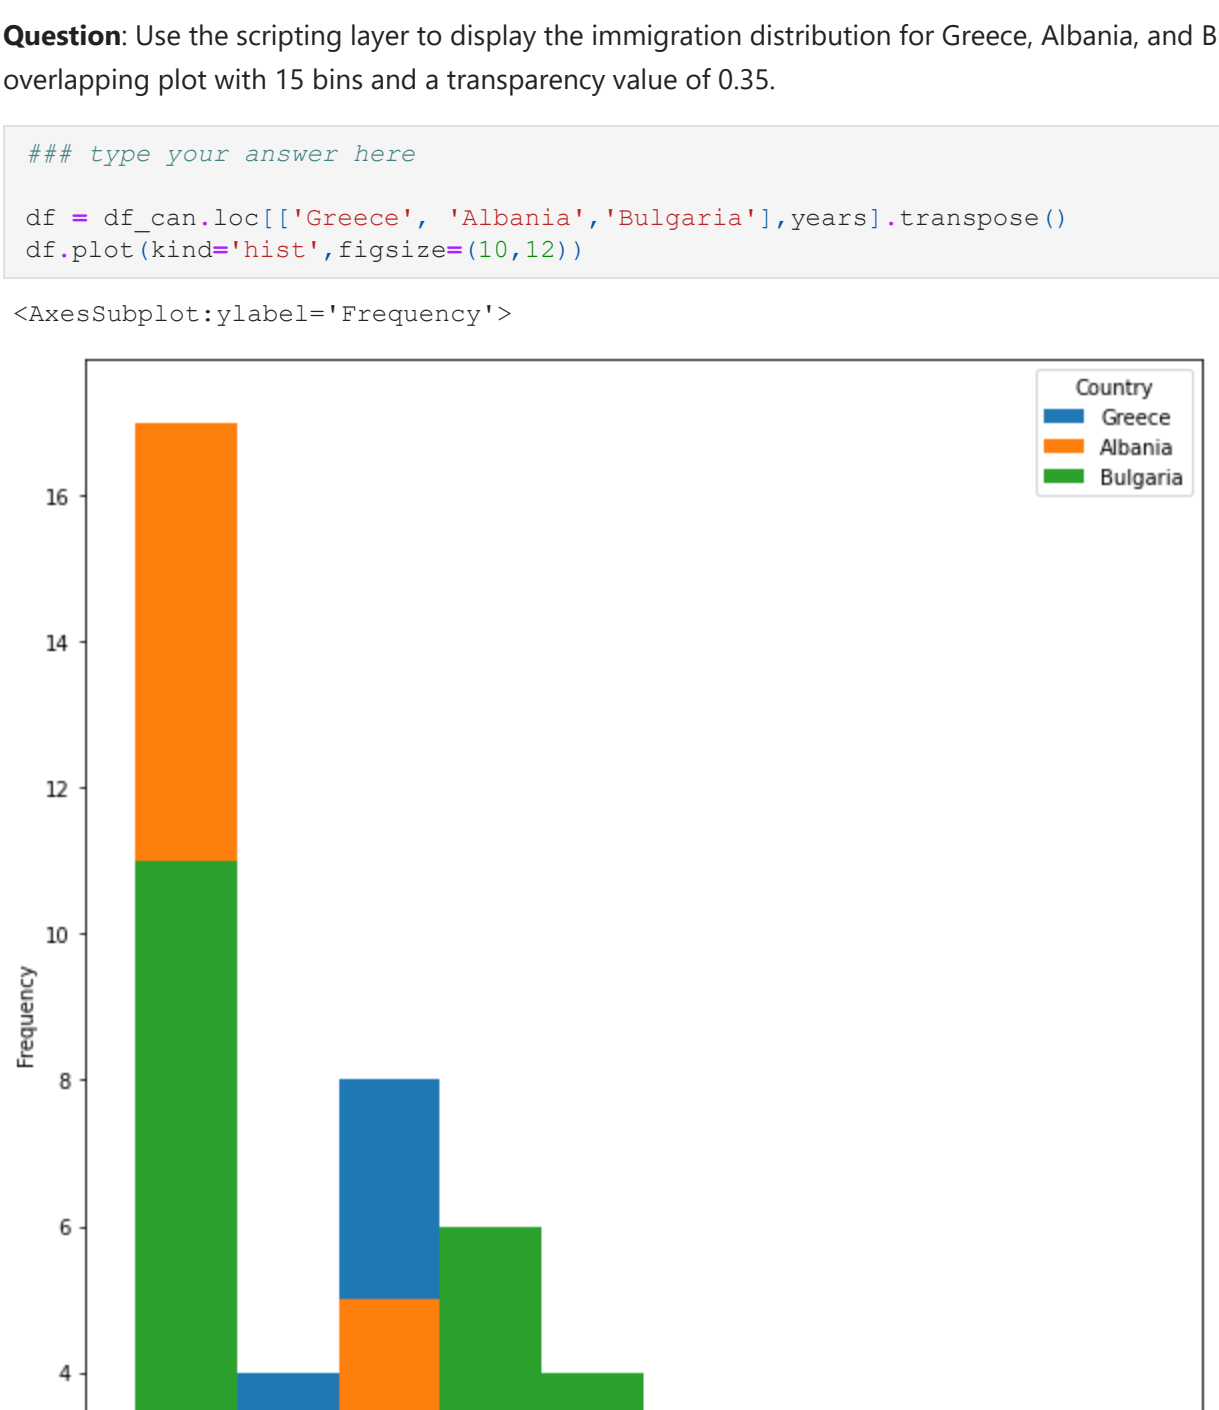
```
import matplotlib
for name, hex in matplotlib.colors.cnames.items():
    print(name, hex)
```

If we do not want the plots to overlap each other, we can stack them using the `stacked` parameter. Let's also adjust the min and max x-axis labels to remove the extra gap on the edges of the plot. We can pass a tuple (min,max) using the `xlim` parameter, as show below.

```
In [44]: count, bin_edges = np.histogram(df_t, 15)
xmin = bin_edges[0] - 10 # first bin value is 31.0, adding buffer of 10 for aesthetic purposes
xmax = bin_edges[-1] + 10 # last bin value is 308.0, adding buffer of 10 for aesthetic purposes

# stacked histogram
df_t.plot(kind='hist',
         figsize=(10, 6),
         bins=15,
         xticks=bin_edges,
         color=['total', 'darkslateblue', 'mediumseagreen'],
         stacked=True,
         xlim=(xmin, xmax))

plt.title('Histogram of Immigration from Denmark, Norway, and Sweden from 1980 - 2013')
plt.xlabel('Number of Years')
plt.ylabel('Number of Immigrants')
plt.show()
```

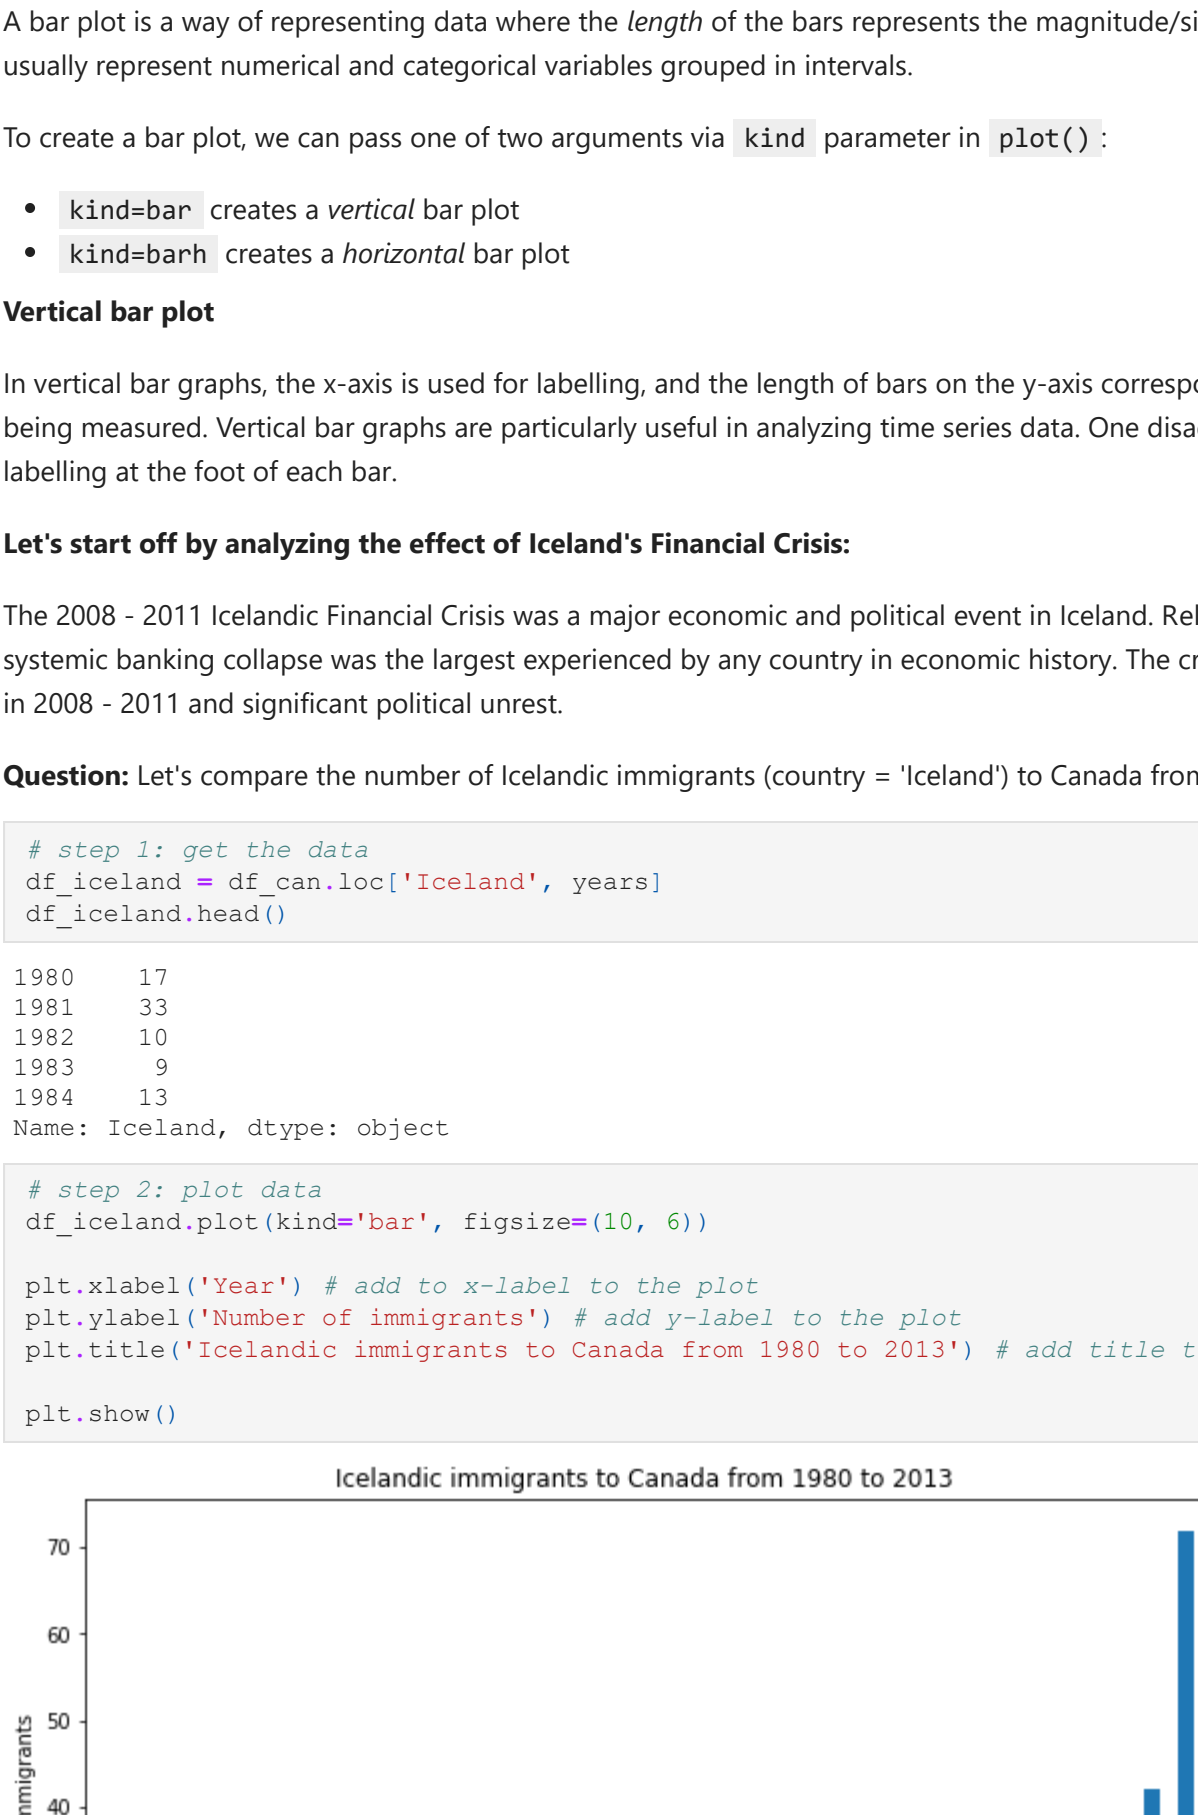


Question: Use the scripting layer with 15 bins and a transparency value for Greece, Albania, and Bulgaria for years 1980 - 2013? Use an overlapping plot with the scripts below to display the value of 0.35.

```
In [54]: ## type your answer here

df = df_can.loc(['Greece', 'Albania', 'Bulgaria'], years)
df.plot(kind='hist', figsize=(10, 12))
```

Out [54]: <AxesSubplot:ylabel='Frequency'>



► Click here for a sample python solution

Bar Charts (Dataframe)

A bar plot is a way of representing data where the *length* of the bars represents the magnitude/size of the feature/variable. Bar graphs usually represent numerical and categorical variables grouped in intervals.

To create a bar plot, we can pass one of two arguments via `kind` parameter in `plot()`:

- `kind=bar` creates a vertical bar plot
- `kind=barh` creates a horizontal bar plot

Vertical bar plot

In vertical bar graphs, the x-axis is used for labelling, and the length of bars on the y-axis corresponds to the magnitude of the variable being measured. Vertical bar graphs are particularly useful in analyzing time series data. One disadvantage is that they lack space for text labelling at the foot of each bar.

Let's start off by analyzing the effect of Iceland's Financial Crisis:

The 2008 - 2011 Icelandic Financial Crisis was a major economic and political event in Iceland. Relative to the size of its economy, Iceland's systemic banking collapse was the largest experienced by any country in economic history. The crisis led to a severe economic depression in 2008 - 2011 and significant political unrest.

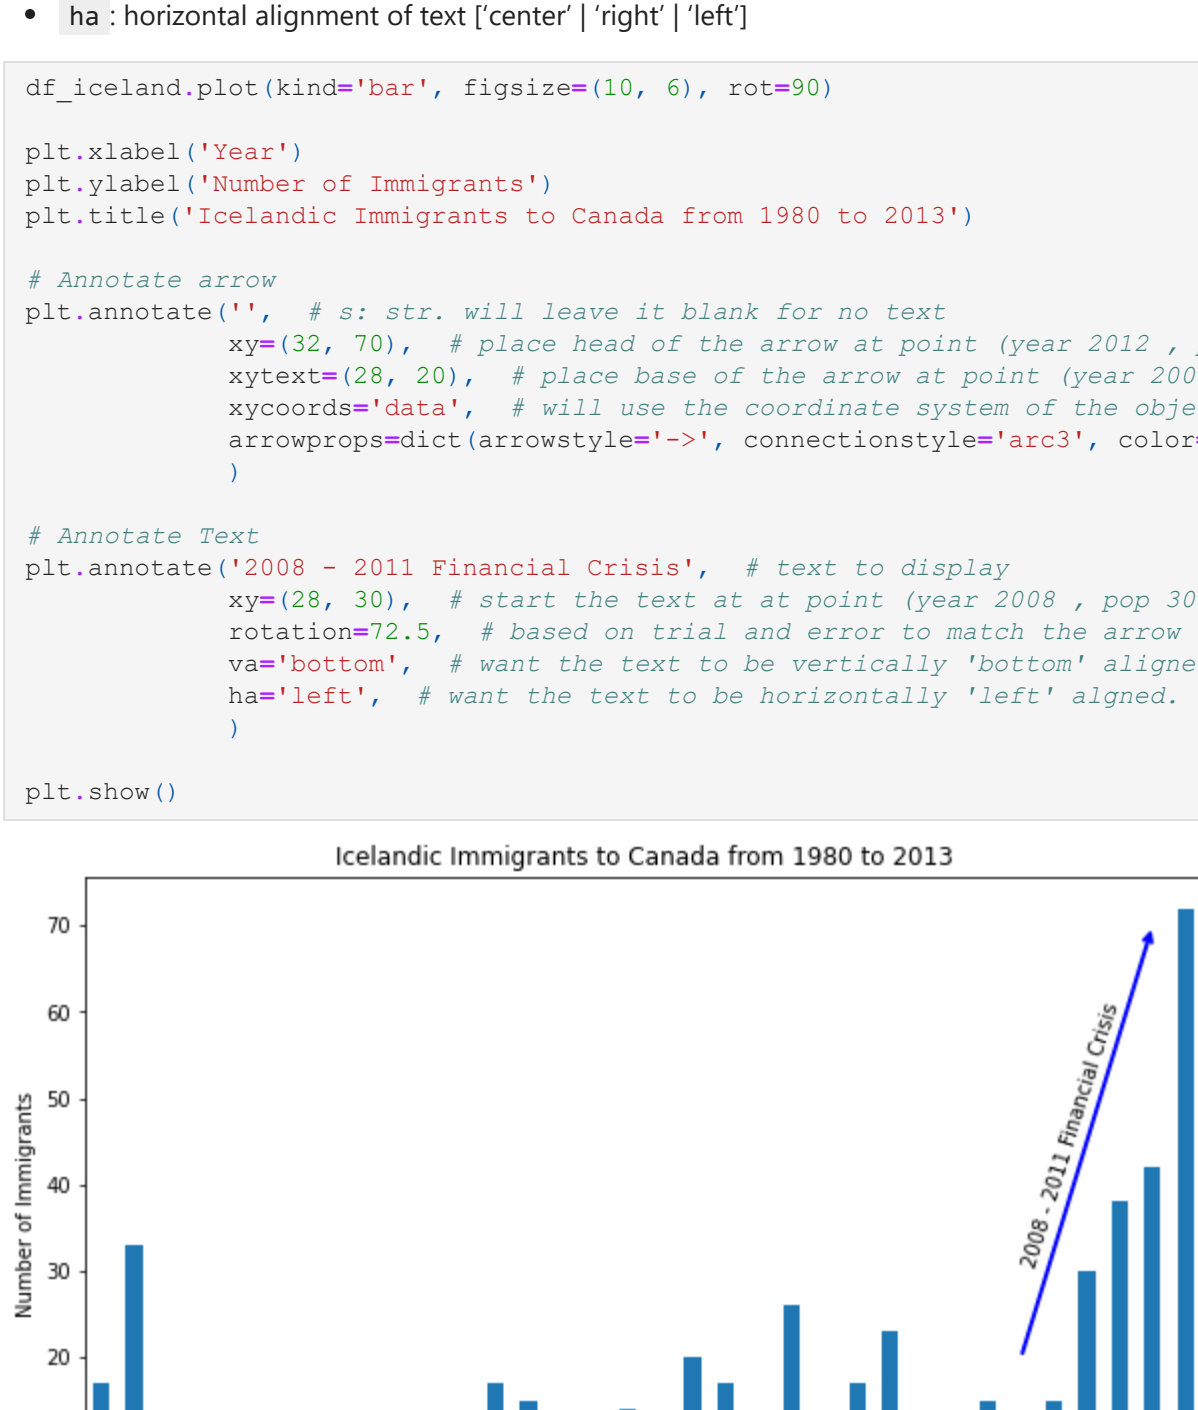
Question: Let's compare the number of Icelandic immigrants (country = 'Iceland') to Canada from year 1980 to 2013.

```
In [55]: # step 1: get the data
df_iceland = df_can.loc['Iceland', years]
df_iceland.head()

Out [55]:
Year    1980    17
        1981    35
        1982    10
        1983     9
        1984    13
        Name: Iceland, dtype: object

In [56]: # step 2: plot data
df_iceland.plot(kind='bar', figsize=(10, 6))

plt.xlabel('Year') # add to x-label to the plot
plt.ylabel('Number of immigrants') # add y-label to the plot
plt.title('Icelandic immigrants to Canada from 1980 to 2013') # add title to the plot
plt.show()
```



The bar plot above shows the total number of immigrants broken down by each year. We can clearly see the impact of the financial crisis: the number of immigrants to Canada started increasing rapidly after 2008.

Let's annotate this on the plot using the `annotate` method of the **scripting layer** or the **pyplot interface**. We will pass in the following parameters:

- `s`: str. the text of annotation.
- `xy`: Tuple specifying the (xy) point to annotate (in this case, end point of arrow).
- `xytext`: Tuple specifying the (xy) point to place the text (in this case, start point of arrow).
- `xycoords`: The coordinate system that xy is given in - 'data' uses the coordinate system of the object being annotated (default).
- `arrowprops`: Takes a dictionary of properties to draw the arrow:
 - `arrowstyle`: Specifies the arrow style. '>' is standard arrow.
 - `connectionstyle`: Specifies the connection type. 'arc3' is a straight line.
 - `color`: Specifies color of arrow.
 - `lw`: Specifies the line width.

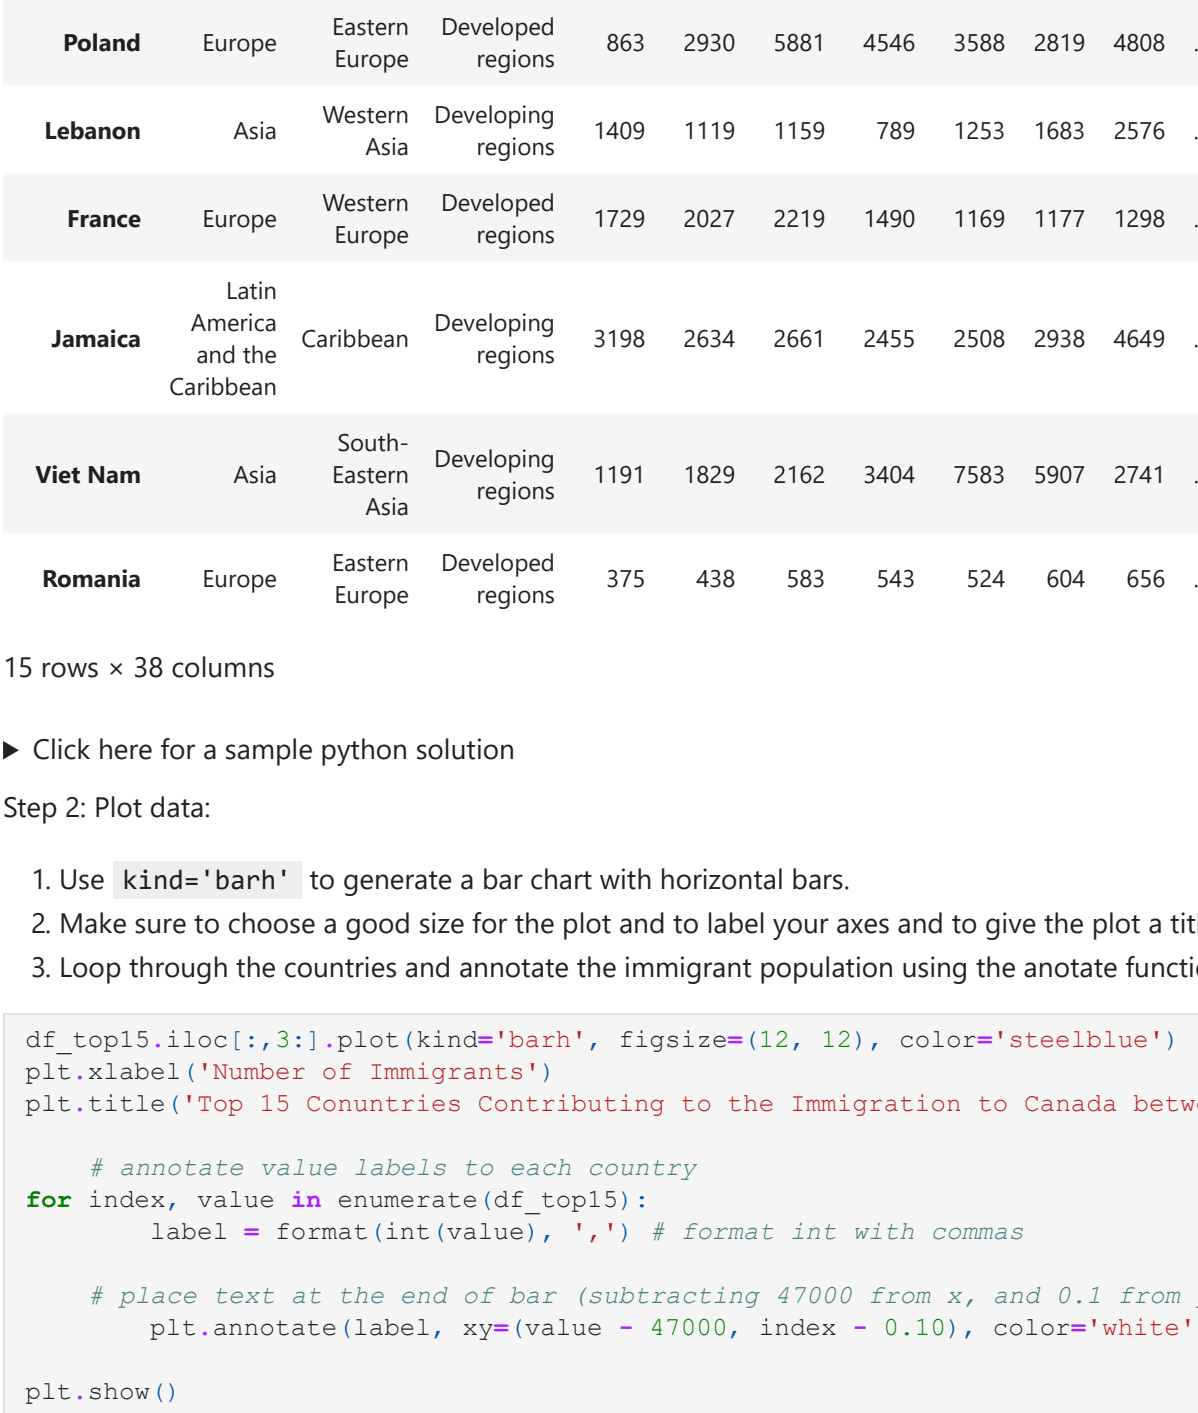
I encourage you to read the Matplotlib documentation for more details on annotations:
http://matplotlib.org/api/pyplot_api.html#matplotlib.pyplot.annotate.

```
In [57]: df_iceland.plot(kind='bar', figsize=(10, 6), rot=90) # rotate the xticks (labelled points on x-axis) by 90 deg

plt.xlabel('Year')
plt.ylabel('Number of immigrants')
plt.title('Icelandic immigrants to Canada from 1980 to 2013')

# Annotate arrow
plt.annotate('', # s: str. Will leave it blank for no text
            xy=(32, 70), # place head of the arrow at point (year 2012, pop 70)
            xytext=(28, 20), # place base of the arrow at point (year 2008, pop 20)
            xycoords='data', # will use the coordinate system of the object being annotated
            arrowprops=dict(arrowstyle='>', connectionstyle='arc3', color='blue', lw=2))

plt.show()
```



Let's also annotate a text to go over the arrow. We will pass in the following additional parameters:

- `rotation`: rotation angle of text in degrees (counter clockwise)
- `va`: vertical alignment of text ['center' | 'top' | 'bottom' | 'baseline']
- `ha`: horizontal alignment of text ['center' | 'right' | 'left']

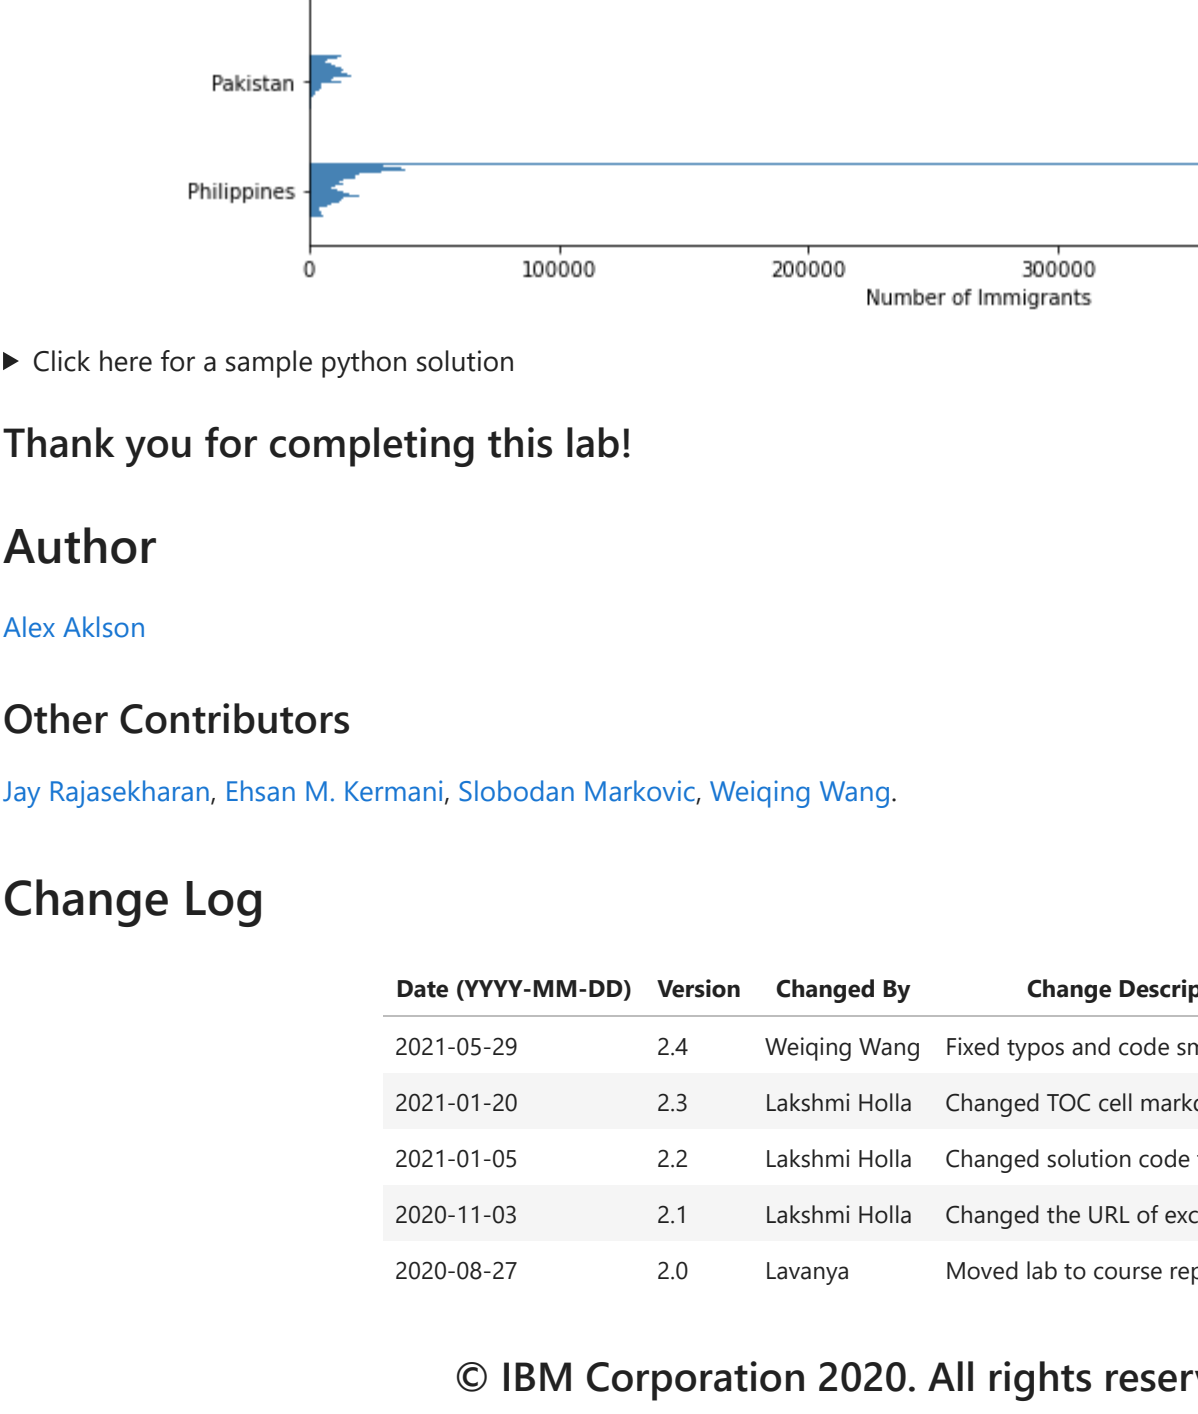
```
In [58]: df_iceland.plot(kind='bar', figsize=(10, 6), rot=90)

plt.xlabel('Year')
plt.ylabel('Number of immigrants')
plt.title('Icelandic immigrants to Canada from 1980 to 2013')

# Annotate arrow
plt.annotate('', # s: str. Will leave it blank for no text
            xy=(32, 70), # place head of the arrow at point (year 2012, pop 70)
            xytext=(28, 20), # place base of the arrow at point (year 2008, pop 20)
            xycoords='data', # will use the coordinate system of the object being annotated
            arrowprops=dict(arrowstyle='>', connectionstyle='arc3', color='blue', lw=2))

# Annotate text
plt.annotate('2008 - 2011 Financial Crisis', # text to display
            xy=(28, 30), # start the text at at point (year 2008, pop 30)
            rotation=2.5, # based on trial and error to match the arrow
            va='bottom', # want the text to be vertically 'bottom' aligned
            ha='left', # want the text to be horizontally 'left' aligned.
            )

plt.show()
```



Horizontal Bar Plot

Sometimes it is more practical to represent the data horizontally, especially if you need more room for labelling the bars. In horizontal bar graphs, the y-axis is used for labelling, and the length of bars on the x-axis corresponds to the magnitude of the variable being measured. As you will see, there is more room on the y-axis to label categorical variables.

Question: Using the scripting layer and the `df_can` dataset, create a horizontal bar plot showing the total number of immigrants to Canada from the top 15 countries, for the period 1980 - 2013. Label each country with the total immigrant count.

Step 1: Get the data pertaining to the top 15 countries.

```
In [75]: df_can.sort_values(by='Total', ascending=False, inplace=True)
df_top15 = df_can.head(15)
```

```
In [76]: df_top15
```

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010
India	Asia	Southern Asia	Developing regions	8880	8670	8147	7338	5704	4211	7150	...	36210	33848	28742	28261	29456	34235
China	Asia	Eastern Asia	Developing regions	5123	6682	6308	1863	1527	1816	1960	...	42584	33518	27642	30037	29622	30391
United Kingdom of Great Britain and Northern Ireland	Europe	Northern Europe	Developed regions	22045	24796	20620	10015	10170	9564	9470	...	7258	7140	8216	8979	8876	8724
Philippines	Asia	South-Eastern Asia	Developing regions	6051	5921	5249	4562	3801	3150	4166	...	18139	18400	19837	24887	28573	38617
Pakistan	Asia	Southern Asia	Developing regions	978	972	1201	900	668	514	691	...	14314	13127	10124	8994	7217	6811
United States of America	Northern America	Northern America	Developed regions	9378	10030	9074	7100	6661	6543	7074	...	8394	9613	9463	10190	8995	8142
Iran (Islamic Republic of)	Asia	Southern Asia	Developing regions	1172	1429	1822	1592	1977	1648	1794	...	5837	7480	6974	6475	6580	7477
Sri Lanka	Asia	Southern Asia	Developing regions	185	371	290	197	1086	845	1838	...	4930	4714	4123	4756	4547	4422
Republic of Korea	Asia	Eastern Asia	Developing regions	1011	1456	1572	1081	847	962	1208	...	5832	6215	5920	7294	5874	5537
Poland	Europe	Eastern Europe	Developed regions	863	2930	5881	4546	3588	2819	4808	...	1405	1263	1235	1267	1013	795
Lebanon	Asia	Western Europe	Developed regions	1409	1119	1159	789	1253	1683	2576	...	3709	3802	3467	3566	3077	3432
France	Europe	Western Europe	Developed regions	2227	2219	1490	1169	1177	1298	...	4429	4002	4290	4532	5051	4646	
Jamaica	Latin America and the Caribbean	Caribbean	Developing regions	3198	2634	2661	2455	2508	2938	4649	...	1945	1722	2141	2334	2456	2321
Viet Nam	Asia	South-Eastern Asia	Developing regions	1191	1829	2162	3404	7583	5907	2741	...	1852	3153	2574	1784	2171	1942
Romania	Europe	Eastern Europe	Developed regions	375	438	583	543	524	604	656	...	5048	4468	3834	2837	2076	1922

15 rows × 38 columns

► Click here for a sample python solution

Step 2: Plot data:

- Use `kind='barh'` to generate a bar chart with horizontal bars.
- Make sure to choose a good size for the plot and to label your axes and to give the plot a title.
- Loop through the countries and annotate the immigrant population using the `annotate` function of the scripting interface.

```
In [78]: df_top15.iloc[:15, 3:].plot(kind='barh', figsize=(12, 12), color='steelblue')
plt.xlabel('Number of immigrants')
plt.title('Top 15 Countries Contributing to the Immigration to Canada between 1980 - 2013')

# annotate value labels to each country
for index, value in enumerate(df_top15):
    label = format(int(value), ',') # format int with commas

# place text at the end of bar (subtracting 47000 from x, and 0.1 from y to make it fit within the bar)
plt.annotate(label, xy=(value - 47000, index - 0.10), color='white')

plt.show()
```



► Click here for a sample python solution

Thank you for completing this lab!

Author

Alex Aikson

Other Contributors

Jay Rajasekharan, Ehsan M. Kermani, Slobodan Markovic, Weiqing Wang.

Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2021-05-29	2.4	Weiqing Wang	Fixed typos and code smells.
2021-01-20	2.3	Lakshmi Holla	Changed TOC cell markdown
2021-01-05	2.2	Lakshmi Holla	Changed solution code for annotate
2020-11-03	2.1	Lakshmi Holla	Changed the URL of excel file
2020-08-27	2.0	Lavanya	Moved lab to course repo in GitLab

© IBM Corporation 2020. All rights reserved.